

3x4 Gridworld Problem with Deep Q-Learning

Kaito Namatame
Arizona State University
Email: knamatam@asu.edu

Code available at: <https://colab.research.google.com/drive/1pOFc29VGXSXg6-mNYUjZcDYGNm28OVfw?authuser=2>
The learning curves can be seen on the https://wandb.ai/gridworld_qlearning/gridworld_dqn_sweep_8

Abstract—This report presents an implementation of Deep Q-Learning (DQN) to solve a 3x4 gridworld task. The agent learns to reach the goal while avoiding a lose state under stochastic transitions. A neural network approximates the Q-function using experience replay and a target network. The training process and learning curves are analyzed, showing that the agent achieves stable convergence and strong average evaluation performance.

Index Terms—Deep Q-Learning, Gridworld

I. INTRODUCTION

In the previous assignment, we implemented tabular Q-learning. In this work, we extend it to Deep Q-Learning (DQN) [1], which uses a neural network to approximate the Q-value function. The 3x4 gridworld is used as a benchmark where the agent learns to reach the goal while avoiding the lose and wall state. We analyze training performance and the effect of hyperparameters using learning curves.

II. METHODOLOGY

A. Environment Setup

The 3x4 gridworld has 3 rows and 4 columns, with a goal state at (3,2) (reward +1), a lose state at (3,1) (reward -1), a wall at (1,1), and a start state at (0,0). Non-terminal states cost a reward of -0.04. Actions (up, down, left, right) with 80% probability for desired directions, with 10% probability each for perpendicular directions. The agent stays in place if it hits a wall or boundary.

B. Data Generation

Training data is generated via experience replay. The agent collects transitions (state, action, reward, next state, done) under an ϵ -greedy policy and stores them in a buffer of 1,000. Randomly sampling minibatches of 64 ensures stable learning by breaking temporal correlations.

C. Neural Network Design

The Q-function is approximated by a fully connected neural network with:

- **Input Layer:** 12 neurons (one-hot encoded state for 3x4 grid).
- **Hidden Layers:** Two layers of 64 neurons each, with ReLU activation.
- **Output Layer:** 4 neurons (Q-values for each action).

D. Target Network

A separate target network, identical in architecture to the policy network, stabilizes training by providing consistent Q-value targets. The target network's weights are updated every 10 steps by copying the policy network's weights.

E. Training Algorithm

The DQN algorithm is shown below

Algorithm 1 DQN with Experience Replay

- 1: Initialize replay buffer \mathcal{D} with capacity 1,000
- 2: Initialize policy network $Q(s, a; \theta)$ and target network $Q(s, a; \theta^-)$
- 3: Initialize Adam optimizer with learning rate 1×10^{-3}
- 4: **for** episode = 1 to 1500 **do**
- 5: Reset environment to start state (0,0)
- 6: Initialize episode reward $R = 0$, steps $t = 0$
- 7: **while** not terminal **do**
- 8: With probability ϵ select a random action a_t
- 9: otherwise select a with maximum Q-value: $a_t = \arg \max_a Q(s_t, a; \theta)$
- 10: Execute a , observe reward r , next state s' , done d
- 11: Store transition (s, a, r, s', d) in \mathcal{D}
- 12: Sample minibatch of N transitions from \mathcal{D}
- 13: Compute target: $y = r + (1 - d) \cdot \gamma \cdot \max_{a'} Q(s', a'; \theta^-)$
- 14: Update $Q(s, a; \theta)$ by minimizing Huber loss:

$$L(\theta) = \frac{1}{N} \sum_{i=1}^N \begin{cases} \frac{1}{2}(y_i - Q(s_i, a_i; \theta))^2 & \text{if } |y_i - Q(s_i, a_i; \theta)| \leq 1 \\ |y_i - Q(s_i, a_i; \theta)| - \frac{1}{2} & \text{otherwise} \end{cases}$$

- 15: Increment t , $R \leftarrow R + r$
 - 16: **if** $t \bmod 10 = 0$ **then**
 - 17: Update target network: $\theta^- \leftarrow \theta$
 - 18: **end if**
 - 19: $s \leftarrow s'$
 - 20: **end while**
 - 21: Decay $\epsilon \leftarrow \max(0.01, \epsilon \cdot 0.999)$
 - 22: **end for**
-

F. Evaluation

The agent is evaluated over 100 episodes with $\epsilon = 0$ (greedy policy). Average rewards are computed and logged to W&B.

G. Hyperparameters

36 combinations of hyperparameters are tested:


- **Learning Rate:** {0.0001, 0.001, 0.1}
- **Discount Factor:** {0.9, 0.99}
- **ϵ Decay:** {0.95, 0.999}
- **Initial ϵ :** {1.0} (fixed)
- **Batch Size:** $N = \{2, 128, 512\}$

The selected hyperparameters are $\alpha = 0.001$, $\gamma=0.99$, ϵ decay=0.999, batch size $N=128$, based on the highest average evaluation reward and policy stability, determined via analysis of learning curves.


III. RESULTS AND DISCUSSION



(a) Training reward vs. (b) Training reward vs. (c) Training reward vs. learning rates (0.01, discount factors (0.9, vs. epsilon decay rates (0.999, 0.95).



(d) Evaluation reward vs. (e) Evaluation reward vs. (f) Evaluation reward vs. learning rates (0.01, 0.001, 0.0001). (0.9, 0.99). (0.999, 0.95).



(g) Training reward vs. (h) Evaluation reward (i) Learned policy with batch sizes (2, 128, vs. batch sizes (2, 128, optimal hyperparameters. 512).

Fig. 1: Comparison of DQN learning curves for different hyperparameters.

Top row (a–c) shows the effect of learning rate (0.01, 0.001, 0.0001), discount factor (0.9, 0.99), and epsilon decay rate (0.999, 0.95) on training reward. Middle row (d–f) shows the same hyperparameters’ effect on evaluation reward. Bottom row (g–i) shows the effect of batch size (2, 128, 512) on training and evaluation reward, and the learned policy with optimal hyperparameters.

- **Learning Rate:** A high learning rate (0.1) caused instability in training, as the Q-values fluctuated significantly

due to large gradient updates. This led to oscillating rewards and a failure to converge as loss is getting bigger. In contrast, a low learning rate (1e-5) resulted in slow learning. The learning rate of 1e-3 achieved the best trade-off between stability and learning speed, yielding the highest evaluation reward.

- **Discount Factor:** From Fig. 1 (c), a lower discount factor ($\gamma=0.9$) led to faster but suboptimal learning, while a higher one ($\gamma=0.99$) caused some initial instability in training due to overemphasis on future rewards. However, the higher discount factor ultimately achieved higher average rewards in the evaluation phase and greater stability compared to the lower discount factor.
- **ϵ Decay:** With ϵ -decay = 0.999, the agent continued exploration for a longer period, resulting in slower improvement in the initial training phase, but achieved a higher average evaluation reward. In contrast, with ϵ -decay = 0.95, exploration decreased rapidly, leading to faster initial learning but premature convergence to a suboptimal policy and lower evaluation rewards. These results suggest that maintaining sufficient exploration, even in a small environment, contributes to more stable and effective policy learning.
- **Batch Size:** A small batch size ($N=2$) increases noise in gradient estimates, leading to unstable updates of the Q-value and lower reward, resulting in lower average evaluation reward, shown in Fig 1 (g). In contrast, a larger batch size provides a better average evaluation reward.
- **Hyperparameter Importance:** The learning rate is the most important hyperparameter, as it directly affects Q-value convergence and stability. A high learning rate (0.1) caused instability, with large gradient updates leading to oscillating loss and poor evaluation rewards. In the small state space, excessive updates amplified errors, making it essential to carefully tune the learning rate for effective policy learning.

IV. CONCLUSION

This work demonstrates that Deep Q-Learning effectively solves the 3x4 gridworld problem, with hyperparameter tuning—particularly the learning rate—being most important for stable convergence in a small, stochastic environment. Future work could extend this approach to larger or continuous environments, such as CartPole, where stability and generalization become more challenging.

REFERENCES

- [1] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A Rusu, Joel Veness, Marc G Bellemare, Alex Graves, Martin Riedmiller, Andreas K Fijelnd, Georg Ostrovski, et al. Human-level control through deep reinforcement learning. *nature*, 518(7540):529–533, 2015.